

Experiment-10

Kubernetes Engine: Qwik Start: Deploy a containerized application to a Kubernetes Engine cluster.

Steps to Deploy a Containerized Application to GKE

Set Up Google Cloud SDK and Kubernetes Tools

1. Install Google Cloud SDK

If you haven't already installed the Google Cloud SDK, follow the instructions here:

Google Cloud SDK Installation

2. Install kubectl

kubectl is the Kubernetes command-line tool used to manage Kubernetes clusters. The Cloud SDK includes kubectl, so if you have the SDK installed, you already have kubectl.

Create a Google Cloud Project

1. Create a new Google Cloud project (if you don't already have one):

- Go to the Google Cloud Console.
- Click on **Select a Project > New Project**.
- Name your project and click **Create**.

2. Set your project in the gcloud CLI:

```
gcloud config set project PROJECT_ID
```

Enable Required APIs

1. Enable Kubernetes Engine API:

```
gcloud services enable container.googleapis.com
```

2. Enable Compute Engine API (if not already enabled):

```
gcloud services enable compute.googleapis.com
```

Create a Kubernetes Cluster

1. Create the Kubernetes Engine cluster:

```
gcloud container clusters create my-cluster --zone us-central1-a
```

Replace my-cluster with your desired cluster name and adjust the zone if necessary.

2. Get the credentials for your cluster: This command configures kubectl to use the cluster you just created.

```
gcloud container clusters get-credentials my-cluster --zone us-central1-a
```

Create a Containerized Application

1. Create a Dockerfile for your application. Below is an example for a simple web application using Node.js:

Dockerfile:

dockerfile

CopyEdit

FROM node:14

WORKDIR /usr/src/app

COPY . .

RUN npm install

EXPOSE 8080

CMD ["npm", "start"]

2. Build the Docker image:

docker build -t gcr.io/PROJECT_ID/my-app:v1 .

Replace PROJECT_ID with your Google Cloud project ID.

3. Push the image to Google Container Registry:

docker push gcr.io/PROJECT_ID/my-app:v1

Create a Kubernetes Deployment

- 1. Create a Kubernetes Deployment configuration file**
(deployment.yaml) for your containerized application.

deployment.yaml:

apiVersion: apps/v1

kind: Deployment

metadata:

```
name: my-app
```

```
spec:
```

```
replicas: 3
```

```
selector:
```

```
  matchLabels:
```

```
    app: my-app
```

```
template:
```

```
  metadata:
```

```
    labels:
```

```
      app: my-app
```

```
  spec:
```

```
    containers:
```

```
      - name: my-app
```

```
        image: gcr.io/PROJECT_ID/my-app:v1
```

```
        ports:
```

```
          - containerPort: 8080
```

Replace PROJECT_ID with your project ID.

2. Apply the Deployment to the Kubernetes cluster:

```
kubectl apply -f deployment.yaml
```

Expose the Application via a Service

1. **Create a Service to expose the application** (either LoadBalancer or ClusterIP for internal access).

Example **Service.yaml** for external exposure (LoadBalancer type):

apiVersion: v1

kind: Service

metadata:

name: my-app-service

spec:

selector:

app: my-app

ports:

- protocol: TCP

port: 80

targetPort: 8080

type: LoadBalancer

2. **Apply the Service configuration:**

kubectl apply -f service.yaml

3. **Get the external IP address:** It may take a few moments for the LoadBalancer to be provisioned.

kubectl get svc

The **EXTERNAL-IP** column will show the public IP once the LoadBalancer is provisioned.

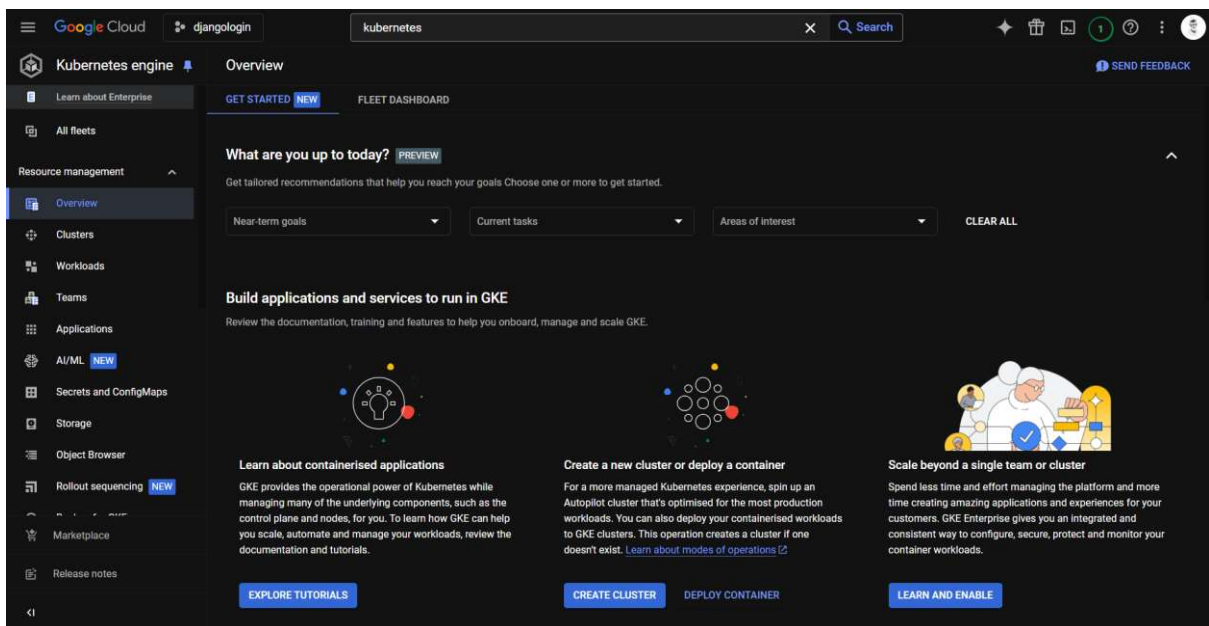
Verify the Application

1. Open a web browser and navigate to the external IP address to verify your application is running.
2. You can also use kubectl to get the status of your pods and services:

```
kubectl get pods
```

```
kubectl get svc
```

Output



Create a deployment

HIDE INFO PANEL

LEARN

1 Deployment configuration

Define how Kubernetes deploys, manages and scales container image

2 Container details

3 Expose (optional)

Define how your deployment is exposed

Deployment configuration

A deployment is a configuration which defines how Kubernetes deploys, manages and scales your container image. Kubernetes will ensure that your system matches this configuration. Three replicas will be created by default.

Deployment name *

deployment-1

Namespace *

default

Labels

Key 1 *

app

Value 1

deployment-1

+ ADD KUBERNETES LABEL

Cluster

New cluster will be in Autopilot mode (see [here](#)). GKE will automatically provision, configure and manage the resources and hardware.

Your deployment will use compute instances managed in a logical grouping called a 'cluster', which will be configured in a way that's great for getting started with Kubernetes.

The cluster will be named deployment-1-cluster

Deployments

Kubernetes Engine works with [containerised applications](#): applications packaged into hardware-independent, isolated user-space instances, for example by using [Docker](#).

In Kubernetes Engine and Kubernetes, these containers are collectively called workloads.

Continuous delivery

Looking for fully managed continuous delivery with multi-environment progressions, scaling and security built-in? [Try Google Cloud Deploy](#)

Recommended for you

Overview of deploying workloads

Help document

Create Kubernetes controller objects to deploy and manage containerized applications and other workloads on a cluster.

Stateless applications

Help document

Deploy a stateless Linux application using GKE.

Deploying a stateful application

Help document

Deploy Stateful applications in GKE, with stored data accessible to the server, clients, and other applications.

Scaling applications

Help document

Scale a deployed application in GKE.

Performing rolling updates

Help document

Perform rolling updates for applications in GKE.

Exposing applications using services

DEPLOY

CANCEL

VIEW YAML